

```

1
2 # 2.1: Introduction to Classification
3
4 + can be infinite number of "equally good" classification thresholds
5 + check: [visuals]
6
7 # 2.2: Choosing a Classifier
8
9 + costs of misclassifying must be considered (i.e. we may be willing to misclassify at
10 a higher rate to reduce the chance of the opposite condition)
11 + check: when the classification plane is horizontal in a 2-d plot, this indicates that
12 the factor plotted on the vertical axis is most important
13 + check: soft classifier must be used when it is impossible to create a perfectly
14 separating hyperplane
15
16 # 2.3: Data Definitions
17
18 + data types... e.g.. data points; attributes vs. features; responses vs. outcomes;
19 structured data vs. unstructured data (e.g. text); types of structured data, including
20 quantitative (numbers with meaning), categorical (numbers without meaning) vs. binary
21 (subset of categorical data), unrelated, time-series
22 + check: [examples]
23
24 # 2.4: Support Vector Machines (SVM)
25
26 + margin formulation
27   + n number of data points,
28   + m attributes,
29   +  $x_i$  = ith attribute of nth data point,
30   +  $y_j$  = response for data point j (1 if data point is positive and -1 otherwise)
31   + line: (sum from  $i = 1$  to  $m$  ( $a_i * x_{ij}$ )) +  $a_0 = 0$ 
32   distance between two lines =  $a / (\text{sqrt}(\text{sum over } i (a_i)^2))$ 
33   + therefore, hard separation objective: minimize across all  $a_j$  (sum from  $i = 1$  to
34  $m$  ( $a_j^2$ )) such that ((sum from  $i = 1$  to  $m$  ( $a_i * x_{ij}$ )) +  $a_0$ ) *  $y_j \geq 1$  for
35 each data point j
36   + objective interpretation = maximize the margin between separating lines OR
37 minimize sum of squares of coefficients
38
39 + error
40   + error =  $\max(0, 1 - ((\text{sum from } i = 1 \text{ to } m (a_i * x_{ij})) + a_0) * y_j)$ 
41   + correct = ((sum from  $i = 1$  to  $m$  ( $a_i * x_{ij}$ )) +  $a_0$ ) *  $y_j \geq 1$  (i.e. see the
42 hard separation objective formulation)
43   + wrong = ((sum from  $i = 1$  to  $m$  ( $a_i * x_{ij}$ )) +  $a_0$ ) *  $y_j - 1 < 0$ 
44   + therefore, total_error = sum from  $j = 1$  to  $n$  (error)
45
46 + final formula
47   + minimize a combination between margin (which we want to maximize) and error
48   (which we want to minimize) minimize across all  $a_j$  (total_error + gamma * sum from
49  $i = 1$  to  $m$  ( $a_j^2$ ))
50 + check: [formula]
51
52 # 2.5: SVM: What the Name Means
53
54 + support vector = points that "hold up" the shape
55 + support vector machine (model) automatically determines support vectors (points
56 supporting shape on parallel lines)
57 + check: [none]
58
59 # 2.6: Advanced SVM
60
61 + can use a multiplier  $m_j$  for error term (in a similar fashion to gamma for the
62 coefficients term) to increase penalty of error
63 + important to scale the data so that minimization of coefficients is robust
64 + near-zero coefficients are probably not relevant
65 + alternatives: kernels for non-linear classifiers; logistic regression for probabilities
66 + check: [formula]
67
68 # 2.7: Scaling and Standardization

```

```

54 + common to scale data to range of 0 to 1;  $x_{ij\_scaled} = (x_{ij} - x_{j\_min}) / (x_{j\_max} - x_{j\_min})$ 
55 + also may standardize (where mean = 0 and sd = 1)  $\rightarrow x_{ij\_standardized} = (x_{ij} - \mu_j) / (\sigma_j)$ 
56 + which method? shorthand: scale for data in a bounded range, standardize for models (e.g. PCA and clustering)
57 + check: [example of scaling]
58
59 # 2.8: K-Nearest Neighbor Classification
60
61 + determine the class of a new point by picking the k closes points to the new one and choosing the class to be the most common among these k "neighbors"
62 + what is nearest? use a p-order distance; may consider adding weighting attributes by importance (which allows for unimportant attributes to be ignored)
63 + how to choose k? different training/validation/testing sets
64 + check: [visual example]
65
66 # 3.1: Introduction to Validation
67
68 + real effect = real relationship between attributes and response
69 + random effect = random, but looks like a real effect
70 + purpose of model validation is to try to distinguish real and random effects
71 + can't measure models' effectiveness on data it is trained on because model fit captures real AND random effects, and only the real effects are likely to manifest in other data
72 + check: if we use the same data to fit a model as we do to estimate how good it is, the model will appear to be better than it really is
73
74 # 3.2: Validation and Test Data Sets
75
76 + split data into training set (larger) to fit model and validation set (smaller) to estimate effectiveness
77 + when comparing two different types of models, need to have a third set (test set), so that training set is used to build and tune models, validation set is used to pick a single model, and test set is used to estimate the performance of the chosen model
78 + if not choosing among more than one model, then test set and validation sets are the same (meaning that there are only two sets total)
79 + check: [basically same as previous lesson]
80
81 # 3.3: Splitting Data
82
83 + rule of thumb: 70-90% training, 10-30% test for one model; 50-70% training and 50-50 split for validation and test sets when comparing models
84 + methods: random; rotation (e.g. 5 data point rotation sequence)
85 + problems: both random and rotation splitting may have issues with time-series data (random may randomly assign an unrepresentative sample to each category, and rotation may introduce bias by allocating too many samples by nature of its sequence)
86 + check: most of the data should be in the training set
87
88 # 3.4: Cross-Valdiation
89
90 + overcomes the issue of "important" data being in only one set
91 + common practice: k = 10
92 + model choice? NONE; do not average coefficients across psplit; should instead train model again using all data for final estimate of coefficients
93 + check: in k-fold cross-validation each part of the data is used k-1 times for training and 1 time for validation
94
95 # 4.1: Introduction to Clustering
96
97 + clustering = grouping
98 + check: [visual clustering]
99
100 # 4.2: Distance Norms
101
102 + n-norm distance (or Minkowski distance)  $n / (n + 1)$  formula:  $(\sum (x_i - y_i)^p)^{1/p}$ 

```

```

103 + check: [2-norm]
104
105 # 4.3: K-means Clustering
106
107 + formulation
108     +  $x_{ij}$  = attribute  $j$  of data point  $i$ 
109     +  $y_{ik} = 1$  if data point  $i$  is in cluster  $k$  and  $0$  if not
110     +  $z_{jk}$  = coordinate  $j$  of cluster center  $k$ 
111     + minimize  $y$  and  $z$  over (sum of  $i$  (sum of  $k$  ( (sum of  $j$  over  $(x_{ij} - z_{jk})^2$ ))  $^$ 
112       (1 / 2))) subject to sum for all  $k$   $y_{ik} = 1$ 
113 + steps (expectation-maximization (EM))
114     0) pick  $k$  cluster centers
115     1) assign each data point to nearest cluster center (centroid)
116     2) recalculate centroids
117     3) repeat steps 1 and 2 until no changes
118 + advantages: fast
119 + check: k-means is a "heuristic" because it isn't guaranteed to get the "best" answer
120
121 # 4.4: Practical Details for K-Means
122
123 + how to improve quality of output?
124     1) run several times, choosing different initial cluster centers
125     2) try different values of  $k$  and pick the number that fits the context
126     3) compare total distance vs  $k$  relationship ("Elbow" diagram) and identify the
127       point at which the marginal benefit of adding another cluster is not worth
128       increasing  $k$ 
129 + check: [elbow diagram]
130
131 # 4.5: Clustering for Prediction
132
133 + assign new point to nearest cluster
134 + use Voronoi diagram to identify area associated with cluster and infer to which
135   cluster a new point would be assigned
136 + check: [visual check]
137
138 # 4.6: Supervised vs. Unsupervised Clustering
139
140 + supervised learning (for clustering) = response is known for each data point;
141   otherwise, it is unsupervised
142
143 + check: [clustering vs. classification]
144
145 # 5.1: Data Preparation
146
147 + scale of the data; outliers...
148
149 # 5.2: Outlier Detection
150
151 + types
152     + point = values are far from the rest (e.g. a point far from others in a scatter
153       plot)
154     + contextual = values isn't far from the rest overall, but is far from points
155       nearby in time (e.g. a large deviation in a time-series, sinusoidal-like curve)
156     + collective = value is missing in a range of points, but it is difficult to
157       identify exactly where (e.g. a "missing spike" in a time-series)
158 + how to detect? box-and-whisker plot for 1-D; fit a model and identify points where
159   model predictions are most inaccurate
160 + check: [outlier graph]
161
162 # 5.3: Dealing with Outliers
163
164 + when to remove? confirmed that outlier is part of "bad data"
165 + with bad data, may consider replacing with imputation
166 + can create a logistic regression model to estimate probability of outliers given
167   certain conditions, and a second model to estimate the value of the replacement data
168   point for the outliers
169 + check: incorrectly-recorded data is a justifiable reason to remove an outlier

```

```

159
160 # 6.1: Introduction to Change Detection
161
162 + check: hypotheses tests are often not sufficient for change detection because they
have high threshold levels, which makes them slow to detect changes
163
164 # 6.2: CUSUM for Change Detection
165
166 + check if  $S_t > T$ , where  $S_t = \max(0, S_{t-1} + (x_t - \mu - C))$ ,  $x_t$  = observed value at
time  $t$ ,  $\mu$  = mean of  $x$ ; when detecting a decrease instead of an increase, then the
positions of  $x_t$  and  $\mu$  switch in the formula
167 + how to choose  $C$  and  $T$ ? depends on trade off of costs of the resolution of the change
to be detected and of signaling for action based on cumulative changes
168 + check: with CUSUM, a higher  $T$  detects changes slower, and is less likely to falsely
detect changes
169
170 # 7.1: Introduction to Exponential Smoothing
171
172 +  $S_t = \alpha * x_t + (1 - \alpha) * S_{t-1}$ 
173 +  $\alpha \rightarrow 0$ : lots of randomness, more weight placed on previous baseline  $S_{t-1}$ 
174 + initial condition:  $S_1 = x_1$ 
175 + check:  $\alpha \rightarrow 1$ : less randomness, more "weight" placed on current observation  $x_t$ 
176
177 # 7.2: Trends and Cyclic Effects
178
179 +  $T_t$ : trend at time period  $t$ 
180 +  $S_t = \alpha * x_t + (1 - \alpha) * (S_{t-1} + T_{t-1})$ 
181 +  $T_t = \beta * (S_t - S_{t-1}) + (1 - \beta) * T_{t-1}$ 
182 + initial condition  $T_1 = 0$ 
183 + cycles: can be additive like trends, or multiplicative
184 + multiplicative seasonality:
185     +  $L$  = length of cycle
186     +  $C_t$  = multiplicative seasonality factor for time  $t$  (inflate/deflate  $x_t$  value)
187     +  $S_t = \alpha * x_t / C_{t-L} + (1 - \alpha) * (S_{t-1} + T_{t-1})$ 
188     +  $C_t = \gamma * (x_t / S_t) + (1 - \gamma) * C_{t-L}$ 
189     +  $C_1, \dots, C_L = 1$ , no initial cyclic effect
190 + check: multiplicative seasonality = seasonal effect is proportional to the baseline
191
192 # 7.3: Exponential Smoothing: What The Name Means
193
194 + check: all past observations are considered when calculating  $S_t$ 
195
196 # 7.4: Forecasting
197
198 + best guess for prediction:  $x_{t+1} = S_t$ , for simple exponential smoothing
199 +  $F_t = \alpha * S_t + (1 - \alpha) * S_t$ , so  $F_{t+1} = S_t$  for simple exponential smoothing
200 +  $F_t$  is calculated similarly for double and triple exponential smoothing (i.e.  $F_t =$ 
 $S_t + T_t$  and  $F_t = (S_t + T_t) * C_{(t+1)-L}$  for multiplicative seasonality where best
estimate of  $C_{(t+1)-L} = C_{t+1}$ )
201 + choose  $\alpha$ ,  $\beta$ , and  $\gamma$  with optimization (i.e.  $\min((F_t - x_t) ^ 2)$ )
202 + check: exponential smoothing is best for short-term forecasting because its forecast
is based primarily on the most recent data points
203
204 # 7.5: ARIMA
205
206 + 3 parts:
207     + Differences: exponential smoothing basic equations; useful if data is stationary
(i.e. mean, variance, etc. are constant over time), but if not stationary, then
need to apply differencing;  $D_1$  = diff of consecutive  $x_t$ ,  $D_2$  = diff of diffs, etc.
208     + Autoregressive: predicting current value based on previous time periods' values
209         + order- $p$  autoregressive model = go back  $p$  time periods; autoregression on
differences = use  $p$  time periods of previous  $x_t$  to predict  $d$ th order diffs
210     + Moving average: using previous errors  $\epsilon_t$  as predictors, where  $\epsilon_t =$ 
 $(\hat{x}_t - x_t)$ 
211         + order- $q$  moving average = go back  $q$  time periods
212     + ARIMA( $n, d, q$ ) looks like order  $D$  with order  $p$  AR with order  $MA$ 
213 + shorthand: ARIMA(0,0,0) = white noise; ARIMA(0,1,0) = random walk; ARIMA( $p, 0, 0$ ) =

```

```

214 AR(p); ARIMA(0,0,q) = MA(q); ARIMA(0,1,1) = exponential smoothing
+ forecasting: ARIMA is better than exponential smoothing when the data is more stable
and/or has less peaks/valleys/outliers
215 + common practice: need 40 past data points
216 + check: [definition of autoregression]
217
218 # 7.6: Generalized Auto Regressive Conditional Heteroscedasticity (GARCH)
219
220 + formula is similar to that of ARIMA, but (1) uses variances and squared errors
instead of observations and linear error returns, and (2) does not use differences of
variances
221 + check: GARCH estimates or forecasts variance
222
223 # 8.1: Introduction to Regression
224
225 + best fit regression line minimizes sum of squared errors, defined by coefficient
estimates a_0, a_1, etc.
226 + check: regression would be used instead of a time series model when there are other
predictors that affect the response (not just previous values of the response variable)
227
228 # 8.2: Maximum Likelihood and Information Criteria
229
230 + likelihood = measure the probability density for any parameter set
231 + Example
232 error ~ N(0, sigma ^ s), iid
233 observations: z_1, ..., z_n
234 model estimates: y_1, ..., y_n
235 MLE: set of parameters that minimize sum of squared errors (i.e. min(sum over i = 1 to n
(z_i - y_i) ^ 2))
236 For regression, where m = number of predictors j, n = number of observations i
237 MLE: min(sum over i = 1 to n (z_i - (a_0 + sum over j = 1 to m (a_j * x_ij)) ^ 2))
238 + AIC = 2 * (m + 1) - 2 ln(L*), where L* = ML value; smaller AIC -> better fit; more
useful if there are lots of points
239 + e ^ ((AIC_1 - AIC_2) / 2) -> % difference in model 1 and 2
240 + BIC = m * ln(n) - 2 * ln(L*); encourages models with fewer predictors; only
applicable when n > m
241 + shorthand: BIC difference > 10 -> smaller BIC model very likely better; 10 > BIC diff
> 6 -> smaller BIC model likely better; 6 > BIC > 2 -> smaller BIC model somewhat
likely to be better
242 + check: simpler models are often better because (1) they are easier to explain, (2) to
understand, and (3) are less likely to be overfit
243
244 # 8.3: Using Regression
245
246 + regression is most useful for descriptive analytics (i.e. understanding relationships
between predictors and the response variable) and predictive analytics, but not so
much for prescriptive analytics (i.e. determining the best course of action)
247 + check: regression is not commonly used for prescriptive analytics
248
249 # 8.4: Causation vs. Correlation
250
251 + check: regression should not be used to determine causation, only to model
relationships
252
253 # 8.5: Transformation and Interactions
254
255 + either response or predictor variables may be transformed in order to create a more
linear relationship
256 + check: a negative sign for an interaction term indicates a decrease in the overall
estimate
257
258 # 8.6: Regression Output
259
260 + p-value
261 + estimates the probability that a coefficient = 0; use 0.05 as a basis for
possible variable elimination, but understand that one might use different
thresholds in order to be more/less conservative

```

262 + gotchas: get small when data set is larger even if strength of relationship between predictor and response is not actually better; only represents a probability, which can always be wrong

263 + other output that are related to p-values: confidence intervals, t-statistic (i.e. coefficient / standard error), coefficient value (which, if low in magnitude when multiplied by the attribute value, sort of "nullifies" the significance indicated by a low p-value), r-squared (i.e. estimate of how much variability model accounts for)

264 + check: the r-squared value must be compared to that of other models to gain insight (i.e. a 0.2 r-squared can be very good, depending on the situation)

265

266 # 9.1: Box-Cox Transformation

267

268 + Used to address assumption of normally distributed data

269 + logarithmic transformation  $t(y) = ((y^{\gamma}) - 1) / \gamma$

270 + check: heteroscedasticity = variance is different in different ranges of the data

271

272 # 9.2: De-Trending

273

274 + trend = increase/decrease of data over time

275 + how? factor-by-factor one-dimensional regression

276 + why? trend could mess up factor-based analysis

277 check: maybe de-trend before using time-series data in a regression model

278

279 # 9.3: Introduction to Principal Component Analysis (PCA)

280

281 + useful for models with lots of dimensions; reduce the amount of data needed (useful if data isn't "complete"); eliminates collinearity; concentrating on only top handful of components can help reduce randomness (because first several components have higher signal-to-noise ratio)

282 + identifies order of dimensions according to amount of "spread" in a given dimensional

283 + check: PCA can eliminate correlation between dimensions and rank dimensions in likely order of importance

284

285 # 9.4: Using PCA

286

287 + how?

288 1) Scale the data to get matrix X: scale such that  $(1 / m) * \sum_i (x_{ij}) = \mu_j = 0$  where  $x_{ij}$  is the  $j$ th factor of data point  $i$

289 2) Find all eigenvectors of  $X^T * X$ , where  $V = [V_1 V_2 \dots]$  is the matrix of eigenvectors

290 3) Find the principal components:  $PC1 = X * V_1$ , etc.;  $k$ th new factor value for  $i$ th data point =  $t_{ik} = \sum_j \text{over } m x_{ij} * v_{jk}$

291 + can have non-linear kernels

292 + interpretation in original factors: more math...

293 + check: in a regression model using PCA the original attributes' implied regression coefficient is a linear combination of the PC's regression coefficients (which is equivalent to the inverse transformation)

294

295 # 9.5: Eigenvectors

296

297 +  $v$  is a vector such that  $A * v = \gamma * v$ , where  $v$  is an eigenvector of  $A$  and  $\gamma$  is an eigenvalue of  $A$  such that  $\det(A - \gamma * I) = 0$

298 + given  $\gamma$ , solve  $A * v = \gamma * v$  to find  $v$

299

300

301 # 10.1: Introduction to Classification and Regression Trees (CART)

302

303 + trees are useful for regression, classification, AND decisions (not necessarily identical to classification)

304 + decision tree: for scenario where separate, inter-dependent decisions need to be made in a process

305 + good if effects of factors are different in different combinations

306 + regression for each leaf in tree

307 + tree approach is useful for pinpointing exactly where a model can get better (via low r-squared values for certain leaves)

308 + check: every leaf has a different model (with different predictors, coefficients, etc., or even different "forms" (i.e. knn, svm, etc.))

```

309 + check: each leaf's model is tailored to its subset of data
310
311 # 10.2: Branching
312
313 + method for branching: most common approach is consider each factor one at a time; for
    example, create a regression model at a given leaf, split on the predictor which would
    lead to the lowest total variance in the splits; stop splitting if the decrease in
    variance is greater than some minimum threshold; then start going backwards and use
    the other split at each pair of splits to identify if error is actually improved by the
    branch; if branch does not improve error, then it is removed (i.e. "pruned")
314 + common practice: stop splitting if leaf would have less than 5% of data
315 + check: don't branch past 5% in order to avoid overfitting
316
317 # 10.3: Random Forests
318
319 + bootstrapped version of trees where many different trees are made (randomly), meaning
    that "weaknesses" are given more weight than they might be given otherwise;
    nonetheless, this can improve the "average" error overall and avoid overfitting; on the
    other hand, explaining the prediction is much more difficult
320 + randomness? one: each tree has different data points, because points are resampled
    WITH replacement; two: when branching, randomly choose subset of all predictors, and
    make choice of predictor from that subset
321 + final tree: average for regression, most common response for classification
322 + check: random forest is not good for interpretation purposes
323
324 # 10.4: Logistic Regression (logistic vs. linear)
325
326 + similarities: transformations of data; interaction terms; variable selection; can be
    used for forests
327 + differences (logistic): longer to calculate; no closed-form solution
328 + model quality: r-squared (fraction of variance explained) for linear; "pseudo"
    r-squared for logistic
329 + thresholding: logistic based on probabilities
330 + ROC curve: 1 - specificity (x-axis) vs. sensitivity (y-axis)
331 + AUC (a.k.a. concordance index) = probability that model give's data point for option
    A a higher response value than for option B; not perfect measure (major limitation is
    that it does not differentiate between cost_fn and cost_fp)
332 + check: logistic can be used when the response is a probability or is binary
333
334 # 10.5: Confusion Matrices
335
336 actual      | model
337 -----
338             | real | spam
339 -----
340     real | tp   | fn
341     spam | fp   | fn
342 + tp, fp, tn, fn definitions...
343 + Guidelines: positive = model says yes; negative = model says no; true: model is
    right; false: model is wrong
344 + Example
345     actual      | model
346     -----
347             | real | spam
348     -----
349     real | 490 | 10
350     spam | 100 | 400
351
352
353 % of email is spam = (100 + 400) / sum(...)
354 % of spam in inbox = 100 / (490 + 100)
355 % of real email lost = 10 / (490 + 10)
356
357 check: tn = email is not spam, model is correct
358
359 # 10.6: Situationally-Driven Comparison

```



```

361 + Evaluating a model's quality: use confusion matrix.
362 + Additionally, need costs
363 total_cost = sum((cnt_tp * cost_tp) + (cnt_fn * cost_fn) + ...)
364
365 + Example
366 same numbers as before;
367 counts imply 50% are spam;
368 counts imply cnt_fn / (cnt_fp + cnt_fn) = (100 / (490 + 100)) = 17% are incorrectly
classified as real;
369 costs: cost_tp = 0, cost_fn = 1, cost_fp = 0.04, cost_fn = 0;
370 so, total_cost = sum(490 * 0 ...) = 14;
371 now, if 40% fp or fn, then
372 total cost = sum((cnt_tp * (6/5) * cost_tp) + (cnt_fn * (4/5) * cost_fn) + ...);
373 so total_cost = sum((490 * (6/5) * 0) ...) = 15.2;
374
375 now, a new model has cnt_fn = 50 and cnt_fp = 50;
376 check: although it's more accurate (i.e. only (50 / (450 + 50)) = 10% are incorrectly
classified as real), it "costs" more
377
378 # 10.7: Advanced Topics in Regression
379
380 + Poisson regression: useful when response follows a Poisson distribution (e.g.
arrivals to airport)
381 + Regression splines (spline = function of polynomials that connect to each other);:
fitted different functions to different parts of the data; multi-adaptive regression
splines (MARS) is one implementation
382 + Bayesian regression: uses Bayes' theorem to update initial estimates; most useful
when there is not much data
383 + k-nearest-neighbor (KNN) regression: useful when there is not estimate of prediction
function; predict response as average response of k closest data points
384 + check: [none]
385
386

```



## # 11.1: Introduction to Variable Selection

- + factor-based models: classification, clustering, regression
- + why limit number of factors?
  1. avoid overfitting: especially problematic when # of factors is => # of data points; causes model to fit to random effects
  2. simplicity:
    - + also implies less data collection is required
    - + reduces chances of identifying insignificant factors
    - + easier to interpret and communicate results (i.e. for prescriptive analysis)
  - + sometimes factors can be illegal to use (e.g. race, gender, religion, etc.), so must be careful not to use corresponding highly correlated factors; also, it can be difficult to prove that an overly complex model avoids using these factors
- + check: simpler models with fewer factors avoids (1) overfitting and (2) difficulty of interpretation (not necessarily low prediction quality nor bias in the most important factors)

## # 11.2: Models for Variable Selection

- + forward selection, backward elimination, and stepwise regression
  - + "greedy": takes action without considering future options
  - + criteria other than p-value can be used to determine factor importance (e.g.  $R^2$ , AIC, BIC)
- + lasso:
  - + adds constraint to standard regression optimization objective of minimizing sum of squared errors by adding a "budget"  $\tau$  to use for sum of coefficients (i.e. adds constraint to minimize size of coefficients, which is similar to SVM)
  - + data must be scaled beforehand
  - + some coefficients may be calculated to be 0
  - + how to choose  $\tau$ ? try different values and
    - + consider number of variables
    - + consider quality of model
- + ridge regression:
  - + similar to lasso, but constrains on a combination of the absolute value of the coefficients and their squares
  - + same issues regarding scaling and choosing parameters (in this case,  $\tau$  and  $\lambda$ )
- + ridge regression:

- + elastic net without the absolute value term
- + no variable selection, but can still lead to better predictive models

+ check: key difference between stepwise and lasso regression?  
lasso requires data to be scaled first

### # 11.3: Choosing a Variable Selection Model

- + comparison:
  - + good for initial analysis: forward selection, backward elimination, stepwise regression; stepwise regression is most common
  - + slower but better: lasso, elastic net
- + regularized regression comparison:
  - + lasso: some coefficients forced to 0 to simplify model (because penalty term is linear)
  - + ridge: coefficients shrink toward 0 (because penalty term is quadratic) to reduce variance in estimate (but adds some bias, which has a tradeoff in prediction error)
- + elastic net:
  - + advantages:
    - + variable selection benefits of lasso
    - + predictive benefits of ridge
  - + disadvantages:
    - + arbitrarily rules out some correlated variables like lasso (e.g. two highly correlated variables might have different "costs", and lasso might choose the one with the higher cost)
    - + underestimates coefficients of very predictive variables like ridge

+ check: when two predictors are highly correlated, ridge (not lasso) regression will usually have non-zero coefficients for both

### # 12.1: Introduction to Design of Experiments (DOE)

- + DOE is useful when we don't have the data and getting a full set of data is either impossible or would take too long
- + examples: different colors in banner ad; retailer's display of "related" products to consumer; "representative" sample of survey respondents
- + comparison and control: e.g. when comparing sales price of cars, need to control for color, age, type of car, etc.
- + blocking factor: a factor that could create variation (e.g. sports car instead of a family car)
- + check: [definition of control]

## # 12.2: A/B Testing

- + examples: collect data regarding click per appearance for banner ads and perform hypothesis test to determine if one is better than the other (can do hypothesis testing "on the fly", (i.e. when significant difference is detected, stop testing alternatives and use better option(s)))
- + a/b testing: choosing between two alternatives
- + a/b testing requirements:
  - + collect data quickly
  - + data must be representative
  - + amount of data is small compared to whole population
- + check: a/b testing is not a good model when collected data is not representative of the population for which we seek insight

## # 12.3 Factorial Designs

- + like a/b testing, but with multiple factors
- + full factorial design example:
  - + test every combination when there are not too many (e.g. 2 fonts x 2 wording x 2 backgrounds)
  - + use ANOVA to determine importance of each factor
- + fractional factorial design:
  - + test subset of combinations when there are many possible combinations (e.g. 7 factors with 3 choices each =  $3^7$  combinations)
  - + create a balanced design where each choice and each pair of choices is tested the same number of times
  - + use regression (possibly including interaction terms, i.e. so as to avoid similar choices for different factors that prove to be the best for their factor, but negatively influence one another when paired) to estimate the effect of each choice
- + check: [an example where factorial design is more appropriate than a/b testing because there are multiple factors]; note that a/b testing is still applicable when there is only one factor with more than 2 choices

## # 12.4: Multi-Armed Bandits

- + useful when you want to continuously test at a rapid pace and want to "maximize" value by also promoting tests that prove to be most "successful"
- + tradeoff of more information vs. immediate value -> "exploration vs. exploitation"

- + exploration: focusing on getting more information (to determine results with more certainty)
- + exploitation: focusing on getting more immediate value
- + theory: from k alternatives...
  - + start with no information, assume equal probability of selecting each alternatives
  - + repeatedly
    - 1) choose an alternative to test based on probability of each being best
    - 2) update probabilities
  - + stop when best alternative is clear
- + parameters
  - + number of tests between recalculating probabilities
  - + method of updating probabilities -> do we assume there is an underlying distribution? if so, which?)
- + there are no real "shorthands", but it is better than running a fixed, large number of tests
- + check: [multi-armed bandit definition]

### # 13.1: Introduction to Advanced Probability Distributions

- + why study these? simple approaches work better sometimes, and probability distribution can form backbone of simple models
- + examples...
- + check: [none]

### # 13.2: Bernoulli, Binomial and Geometric Distributions

- + binomial examples: probability (constant) of people sending donations when charity asks donations from 1/12 of mailing list each month
- + geometric examples: number of interview until first job; number of hits until baseball bat breaks; number of good manufactured units before a defective one
- + geometric (and binomial) assumption: each Bernoulli trial is iid
- + can infer whether data is iid by comparing it to geometric distribution
- + check: binomial distribution is not a good model when estimating the number of days n in each month in which temperature is above a threshold (with probability p) because the results are not independent (e.g. days above threshold are likely to be "clumped" in the summer)
- + check: [geometric formula application]

### # 13.3: Poisson, Exponential and Weibull Distributions

+ weibull extra info:

- +  $k < 1$  -> when failure rate decreases (i.e. "worst things fail fast", like parts with defects)

- +  $k > 1$  -> when failure rate increases (i.e. "things that wear out, like tires)

- + if  $k = 1$  -> Weibull = exponential, where  $\lambda = 1 / \lambda$

+ using software to determine if data fits a probability distribution:

- + input: set of data

- + output: fit of varying distributions and parameters

- + should be used cautiously (e.g. if software finds that distribution is Weibull with  $k = 1.002$ , then it might be better to use exponential with  $k = 1$ )

+ check: if the number of arrivals follows the exponential distribution, then the number of arrivals per unit time follows the Poisson distribution (and visa versa)

+ check: geometric distribution models how many tries it takes for something to happen, while the Weibull distribution models how long it takes

### # 13.4: Q-Q Plots

+ usefulness?

- + for visualizing whether two distributions are about the same

- + for visualizing whether a data set is distributed similarly to a probability distribution

- + note that statistical tests can hide details, so visualization can be better (even if it is not really more quantitative)

+ when comparing a single data set to a probability distribution...

- + horizontal axis = data, vertical axis = theoretical values of percentiles

+ check: [visual example]

### # 13.5: Queuing

+ example description:

- + autodialer automatically calls phone numbers

- + if the call is answered it is put into a queue

- + how many employees should we have?

- + based on how many people will answer the autodialer

- + based on duration of call once the employee is on the phone
- + calls are answered and added by a probability distribution
- + we have c number of employees
- + calls leave the system based on another probability distribution
- + example, simple:
  - + call start is Poisson ( $\lambda$ )
  - + 1 employee
  - + call end is Exponential ( $u$ ) time
  - + we can calculate:
    - + expected fraction of time employee is busy
    - + expected waiting time before talking to employee
    - + expected number of calls waiting in queue
  - + resulting equations:
    - + arrival Rate (calls) =  $\lambda$
    - + service Rate (calls) =  $u > \lambda$
    - + transition Equations ( $\geq 1$  calls in queue)
      - +  $P(\text{next event is an arrival}) = \lambda / (\lambda + u)$
      - +  $P(\text{next event is a finished call}) = u / (\lambda + u)$
    - + can calculate:
      - + Expected fraction of time employee is busy =  $\lambda / u$
      - + Expected waiting time before talking to employee =  $\lambda / u(u + \lambda)$
      - + Expected number of calls waiting in queue =  $\lambda^2 / (u(u + \lambda))$
- + example, more complex: more employees
  - + all can be solved with closed form answers due to memoryless property of the Poisson and Exponential probability distributions
  - + memoryless exponential: distribution of remaining call time = initial distribution of call time
  - + memoryless Poisson: distribution of time to next arrival = initial distribution of time to next arrival
  - + if data fits exponential/Poisson distribution, then it is memoryless, and visa versa
- + memoryless property example :
  - + setting: should tire manufacturer pay damage for accident that happened at 10K miles?
  - + probability (tire fails at 10K) = ?
    - + tires are more likely to fail the more worn out they are, so this is not memoryless
    - + cannot be modeled with the exponential distribution (possibly try the Weibull with  $k > 1$ )

- + potential queuing model parameters:
  - + general arrival distribution (A)
  - + general service distribution (S)
  - + Number of servers (C)
  - + size of the queue (K)
  - + population size (N)
  - + queuing discipline (D)
- + kendall notation
- + model extensions: potential "hang-ups", balking (i.e. leaving after seeing wait time), etc.
- + simulation is good for modeling complex scenerios
- + check: queuing is not appropriate for estimating something not having to do with waiting in line

## # 13.6: Simulation Basics

- + simulation? build a model and watch its behavior
- + types of simulation:
  - + deterministic = same inputs give the same outputs (no randomness)
  - + stochastic = use when system has randomness
- + continuous time simulations: changes happen continuously
- + discrete-event simulations: changes happen at discrete time points only
  - + valuable when systems have high variability
  - + using average values is not good enough
- + simulations software includes:
  - + modeling elements:
    - + entities: things that move through a simulation
    - + modules: parts of process (e.g. queues)
    - + actions
    - + resources (e.g. workers)
    - + decision points
    - + statistical tracking
  - + GUI
  - + complexities "under the hood" (e.g. pseudo random number generation)
- + replications: number of runs of simulation
  - + one replication = one data point (may be unrepresentative)
  - + run multiple times to get distribution of outcomes
- + simulation validation with real data
  - + real and simulated averages don't match -> problem
  - + averages match, variances don't match -> problem



- + check: a stochastic simulations should be run many times because one random outcome might not be representative of system performance in the range of different situations that could arise
- + check: it is important to validate a simulations by comparing it to real data as much as possible because if the simulation isn't a good reflection of reality, then any insights we gain from studying the simulation might not be applicable in reality

### # 13.7: Prescriptive Simulation

- + prescriptive analytics describes how we use simulation for analytics
- + use automated heuristic optimization offered by simulation software
- + when making comparisons, need to be careful
  - + a simple comparison of means may not be sufficient
  - + if possible, compare performance (of different parameter values) on exactly same data point (i.e. use the same random numbers before each trial)
- + simulation can be a powerful tool:
  - + model is only as good as quality of input
  - + missing or incorrect information may lead to incorrect answers
- + example:
  - + in a call center simulation, assuming that workers answer calls equally quickly (which is likely incorrect) can lead to costly bad decisions
- + check: in simulation, both (1) use automated optimization functions in simulation software to find good parameter values, and (2) vary parameters manually

### # 13.8: Markov Chains (MCs)

- + markov chains = probability based models based on states of a system
- + formulation:
  - + for each state  $i$  in the model:
    - +  $p_{ij}$  = transition probability from state  $i$  to state  $j$
    - +  $p = [p_{ij}]$  is the transition matrix
- + example:  $p(\text{sunny} \rightarrow \text{rainy})$  = probability sunny today rainy tomorrow
  - + what is long run of probability of rainy days?
  - + given  $\pi = [0.5, 0.25, 0.25]$  = probabilities of xyz today

- +  $\pi * P$  = probabilities of xyz tomorrow
- +  $\pi * P * P$  = probabilities of xyz the day after tomorrow
- +  $\pi_{\text{star}}$  = "steady state" solution
  - + doesn't always exist
  - + can't have cyclic behavior
  - + every state must be reachable from all others
- + key assumption: memoryless
  - + state transitions only depend on the most recent state
  - + most systems do not exhibit this property though, but we still study them because they are useful in trying to connect smaller amounts of information to find larger ones
- + example: Google's page ranks system
  - + more complex than a simple MC system, but MC is useful for understanding (i.e. we pages = states, chain = jumping from one page to another, steady state probability = rank of web pages)
- + other examples:
  - + urban sprawl, population dynamics, disease propagation
  - + only memoryless in the short term, but that may be all that is needed
- + summary:
  - + no so common due limiting assumption of memoryless system
  - + nonetheless, powerful in some cases
- + check: [memoryless definition] (i.e. memoryless = next state of a process doesn't know much about the form of the underlying distribution the data comes from, or it doesn't an edge between them then remove it; and if there's previous choice beats")
- + information levels:
  - + perfect: know all information (e.g. chess)
  - + imperfect: some may have more information than others (i.e. not symmetric) (e.g. example in next video)
- + zero-sum games:
  - + zero-sum: whatever one side gets, the other side loses (e.g. rock-paper-scissors)
  - + non-zero-sum: total benefit might be higher or lower (e.g. economics)
- + summary:
  - + how to determine best strategy? optimization
- + check: [example] a game theoretic model is appropriate when ... and it must model something as a function of the number of units its competitor produces

## # 16.5a: Competitive Models Demo

- + setup:
  - + two gas stations: bp vs. shell

- + can set price at either \$2.50 or \$2
  - + if both set price equally, then demand is 50/50
  - + otherwise, all demand goes to the lower-priced product
- + scenario 1: cost = \$1/gallon
  - + from shell's POV, best strategy is to sell at \$2 no matter what bp does
  - + lower prices is better for bp as well
- + stable equilibrium:
  - + neither station has incentive to change
    - + "prisoner's dilemma"
      - + even if both sides agree to higher prices, both have incentive to cheat and lower price, so they do that
- + scenario 2: cost = \$1.75/gallon
  - + both are better off charging higher price
- + extension: consider arbitrary prices pshell and pbp
  - + both keep lowering prices until price is about equal to the cost
- + check: [none]

## # 18.1: Introduction to Power Company Case

- + context: power company want to shutoff powr fo customers who don't pay their bills
- + turn power off:
  - + turn off for those not ever going to pay
  - + not people who forgot or got behind
- + logisitical problems:
  - + manually shut off
  - + go to location
  - + more work than the company can handle
- + considerations:
  - + which shutoff should be done?
  - + some worker's time is taken up by travel
  - + how to identify "good" customers whose power should not be shutoff
  - + prioritizing shutoffs

## # 18.2 Models for Customer Identification

- + power customer identification:
  - + customers who can pay, but aren't going to:
    - + credit score
    - + income
    - + past history of defaults on paymebts to any company
    - + past power-bill paymenbt history
    - + sip code
    - + value of home
    - + rent or own
    - + length of residency
    - + marital status
    - + number of residents
  - + some factors may be illegal to use
    - + race, sex, age, other demographic factors, or other factors highly correlated with these
- + types of models:
  - + yes/no answer:
    - + classification (svm or knn):
      - + pay, can pay but do not pay, not able to pay
    - + clustering
  - + probability of payment: e.g. logisitic regression
  - + single-model apporach or treee-based approaches
  - + hybrid approaches: e.g. first cluster, then analyze each cluster separately
- + pros and cons of models:
  - + unsupervised approach:
    - + clustering
      - + quick
      - + not exact cluster you might expect
  - + supervised approach:
    - + classification:
      - + clear decision
    - + logisitic regression:
      - + requires threshold
- + summary:
  - + modeling is an art
  - + not always a "right" approach, but there are wrong approaches

## # 18.3 Models for Cost Estimation

- + power cost estimation: cost of leaving power on or off
  - + for customers with long history...
  - + given customer credit, financial, andpayment history data; and

- possibly some demographic information
  - + use exponential smoothing or ARIMA OR
  - + to estimate the amount of power a customer will use in the next month
- + when considering variability in usage
  - + given [same]
  - + use GARCH
  - + to estimate the amount of variability in a customer's power usage next month
- + OR... for customers with shorter or longer history
  - + given [same]
  - + regression-based model (simple regression, tree-based, or clustering followed by regression)
  - + to [same]
- + pros and cons of models:
  - + time series:
    - + good if enough past customer usage data exists
    - + effective only for short-term forecasts
  - + factor-based regression:
    - + can be effective even when there's not much specific customer data
    - + normalize to account for seasonality
- + hybrid approach: model payment and cost together
  - + model the amount of money owed (zero if bill is paid)
  - + however, this is not usually effective if there are many zero values, plus a range of others, so usually better to analyze separately

#### # 18.4 Models for Shutoff Selection

- + expected cost of leaving power on/off
  - +  $E[\text{cost of keeping power on}] = p(\text{no pay}) * E[\text{cost of power used next month}]$
- +  $(1 - p(\text{no pay})) * 0$ 
  - +  $E[\text{cost of shutting power off}] = p(\text{no pay}) * \text{cost\_shutoff} + (1 - p(\text{no pay})) * (\text{cost\_shutoff} + \text{cost\_turnon})$
- + data
  - + past data on travel time/speed
    - + other details of driving may be too complex/expensive to collect
  - + time to shutoff power
  - + estimate of future usage
  - + probability of non-payment
- + optimization models
  - + question: highest-value set of customers for power shutoffs?
  - + binary variable for each customer
    - + objective function: cost difference between shutting off power or not, times the binary variable
  - + constraints are hard to write
  - + clustering
    - + cluster the physical locations
    - + modify optimization model
  - + simulation
    - + variability in drive times, shut off times, power usage
    - + distribution fitting
  - + then use model to determine how many new workers to hire
- + actual models used
  - + customer identification: logistic regression
  - + cost estimation: linear regression with Box-Cox transformation
  - + shutoff selection: vehicle routing (optimization)